

EE360C: Algorithms

Greedy Algorithms

Part 2/4

Summer 2019

Department of Electrical and Computer Engineering
University of Texas at Austin


Scheduling to Minimize Lateness

Scheduling to Minimize Lateness

Minimizing Lateness Problem

- Single resource processes one job at a time
- Job j requires t_j units of processing time and is due at time d_j
- If j starts at time s_j , it finishes at time $f_j = s_j + t_j$
- Lateness = $l_j = \max\{0, f_j - d_j\}$
- Goal: Schedule all jobs to minimize maximum lateness
 $L = \max l_j$

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15



Greedy Template: Greedy Choice

Greedy Choice

We have to consider the jobs in some order:

- **Shortest processing time first.** Consider jobs in ascending order of processing time t_j .
- **Earliest deadline first.** Consider jobs in ascending order of deadline d_j .
- **Smallest slack.** Consider jobs in ascending order of slack $d_j - t_j$.

Can you construct counter examples for any of these options?

Greedy Algorithm

Sort n jobs by deadline so that $d_1 \leq d_2 \leq \dots \leq d_n$

$t \leftarrow 0$


for $j = 1$ to n

Assign job j to interval $[t, t + t_j]$

$s_j \leftarrow t, f_j \leftarrow t + t_j$

$t \leftarrow t + t_j$


output intervals $[s_j, f_j]$



No Idle Time

Observation 1

There exists an optimal schedule with **no idle time**.



Observation 2

The greedy schedule has no idle time.

Process

Consider an optimal schedule O . Gradually modify O , preserving its optimality at each step, but eventually transforming it into a schedule A that our greedy algorithm would give us. This is an **exchange argument**.

Inversions

Definition

An **inversion** in schedule S is a pair of jobs i and j such that $d_i < d_j$ but j scheduled before i .



Observation

A greedy schedule has no inversions.

Inversions (cont.)

Claim

All schedules with no idle time and no inversions have the same lateness.

Proof



- If two different schedules have no inversions or idle time, then they differ in the order in which jobs that have identical deadlines can be scheduled.
- Let such a deadline be d .
- Jobs with deadline d are scheduled consecutively after jobs with earlier deadlines and after jobs with later deadlines.
- Among jobs with deadline d only the last one has the greatest lateness, and this lateness does not depend on the order of the jobs.

Exchange Argument

Claim

There is an optimal solution that has no inversions and no idle time.

Proof

- Already claimed there is an optimal solution with no idle time.
- (a) If \mathcal{O} has an inversion, then there is a pair of jobs i and j such that j is scheduled immediately after i and has $d_j < d_i$.
 - Consider an inversion in which job a is scheduled sometime before a job b , and $d_a > d_b$.
 - If we advance in the scheduled order of jobs from a to b one at a time, there has to come a point where the deadline we see decreases for the first time.
 - This corresponds to a pair of consecutive jobs that form an inversion.

Exchange Argument (contd.)

Proof

- (b) After swapping i and j we end up with a schedule with one less inversion.
- When we swap i and j the inversion is fixed and no new inversions are created.
- (c) The new swapped schedule has a maximum lateness no larger than that of \mathcal{O} .
- Proof on white board.



Analysis of Greedy Algorithm

Claim

The schedule A produced by the greedy algorithm has optimal maximum lateness L .

Proof

- By last claim, there is an optimal schedule with no inversions and idle time.
- All schedules with no inversions and no idle time have the same lateness.
- Therefore the greedy algorithm schedule A is optimal.

Recap: Greedy Analysis Strategies

Greedy algorithm stays ahead

Show that after each step of the greedy algorithm, its solution is at least as good as any other algorithm's.

Exchange argument

Gradually transform any solution to the one found by the greedy algorithm without hurting its quality.

Structural

Discover a simple “structural” bound asserting that every possible optimal solution must have a certain value. Then show that your algorithm always achieves this bound.

Shortest Paths

Shortest Paths

Given a graph and its vertices and edges, how do we find the shortest path from one vertex to another?

- Clearly, enumerating all of the possible paths, summing the weights of the edges, and taking the minimum is excessively expensive.

This is another pervasive problem:

- networking algorithms
- maps and distances traveled
- scheduling algorithms
- or anything else where the weights can be costs, time, distances, etc.

Shortest Path Problem

The shortest-paths problem

Given a weighted, directed graph $G = (V, E)$ with weight function $w : E \rightarrow \mathbf{R}$, define the weight of a path $p = \langle v_0, v_1, \dots, v_k \rangle$ to be the sum of the weights of its constituent edges:

$$w(p) = \sum_{i=1}^k w(v_{i-1}, v_i)$$

We define the **shortest path weight** from u to v as:

$$\delta(u, v) = \begin{cases} \min\{w(p) : u \overset{p}{\rightsquigarrow} v\} & \text{if there is a path from } u \text{ to } v \\ \infty & \text{otherwise} \end{cases}$$

A **shortest path** from u to v is any path p with weight $w(p) = \delta(u, v)$.

Variants of the Shortest Path Problem

We'll focus first on the **single-source shortest-paths problem** (i.e., find the shortest paths given a specific source vertex), but others exist:

- **single-destination shortest-paths problem**: find the shortest paths to a given destination from all other vertices
- **single-pair shortest-path problem**: given u and v , find the shortest path from u to v
- **all-pairs shortest-paths problem**: find a shortest path for every pair of vertices

Shortest Paths Substructure

The algorithms we'll discuss often rely on the fact that a shortest path contains within it other shortest paths.

Lemma 24.1

Given a weighted, directed graph $G = (V, E)$ with weight function $w : E \rightarrow \mathbf{R}$, let $p = \langle v_0, v_1, \dots, v_k \rangle$ be a shortest path from v_0 to v_k and, for any i and j such that $1 \leq i \leq j \leq k$, let $p_{ij} = \langle v_i, v_{i+1}, \dots, v_j \rangle$ be the subpath of p from v_i to v_j . Then p_{ij} is a shortest path from v_i to v_j .

Prove this by contradiction.

Gotcha: Negative-Weight Edges

At times, it makes sense for edges to have negative weights.

- if there are no cycles in the graph that contain negative weights, then the shortest path weights are well-defined
- if there are cycles that contain negative edges, the shortest path is not well-defined (since if I found one, I could always find one smaller by going around again)
- if there is such a negative edge cycle on a path from u to v , we define $\delta(u, v) = -\infty$

Some algorithms assume that all edge weights are non-negative (so they're basically preceded by a phase that checks this); others can handle negative weights.

Representing Shortest Paths

Usually when we construct the shortest paths, we want more than their weights, we want to know the actual paths, too. We maintain the **predecessor subgraph** $G_\pi = (V_\pi, E_\pi)$. V_π is the set of vertices in G that end up having non-NIL predecessors (and the source, s):

$$V_\pi = \{v \in V : v.\pi \neq \text{nil} \cup s\}$$

E_π is the set of edges induced by the π values for the elements of V_π :

$$E_\pi = \{(v.\pi, v) \in E : v \in V_\pi - \{s\}\}$$

Practically, we maintain $v.\pi$ for each v and create a chain of predecessors that runs back along the shortest path to s . Conceptually, π gives us a tree rooted at s that is a subgraph of G , contains all vertices reachable from s with their shortest paths.

Initialization

For each vertex, we maintain $v.d$ to be an upper estimate on the weight of the shortest path from s to v .

INITIALIZE-SINGLE-SOURCE(G, s)

```
1  for each vertex  $v \in G.V$ 
2      do  $v.d \leftarrow \infty$ 
3       $v.\pi \leftarrow \text{NIL}$ 
4   $s.d \leftarrow 0$ 
```

Relaxation

Relaxation of an edge (u, v) tests whether we can improve the shortest path known to v by going through u , and, if so, updating $v.d$ and $v.\pi$.

RELAX(u, v, w)

```
1  if  $v.d > u.d + w(u, v)$ 
2      then  $v.d \leftarrow u.d + w(u, v)$ 
3           $v.\pi \leftarrow u$ 
```

Properties of Shortest Paths

- **Triangle Inequality:** For any edge $(u, v) \in E$, we have $\delta(s, v) \leq \delta(s, u) + w(u, v)$
- **Upper-bound property:** We always have $v.d \geq \delta(s, v)$ for all $v \in V$, and once $v.d$ reaches $\delta(s, v)$, it never changes
- **No-path property:** If there is no path from s to v , then we always have $v.d = \delta(s, v) = \infty$
- **Convergence property:** If $s \rightsquigarrow u \rightarrow v$ is a shortest path, and if $u.d = \delta(s, u)$ prior to relaxing (u, v) , then $v.d = \delta(s, v)$ after
- **Path-relaxation property:** If $p = \langle v_0, v_1, \dots, v_k \rangle$ is a shortest path from $v_0 = s$ to v_k and the edges of p are relaxed in order, then $v_k.d = \delta(s, v_k)$, regardless of what other relaxations are performed, even if they are interleaved.
- **Predecessor-subgraph property:** Once $v.d = \delta(s, v)$ for all $v \in V$, the predecessor subgraph is a shortest paths tree rooted at s .


Dijkstra's Algorithm

The Intuition

Dijkstra's algorithm maintains a set S of vertices to which the shortest path has been determined.

- the algorithm selects a new vertex u with the minimum shortest path estimate of those left in $V - S$ to add to S
- then relax all of the edges leaving u and repeat
- we use a minimum priority queue Q of the vertices that are keyed by the d estimates

Dijkstra's Algorithm Example



Dijkstra's Algorithm (cont.)

DIJKSTRA(G, w, s)

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S \leftarrow \emptyset$ 
3   $Q \leftarrow G.V$ 
4  while  $Q \neq \emptyset$ 
5      do  $u \leftarrow \text{EXTRACT-MIN}(Q)$ 
6       $S \leftarrow S \cup \{u\}$ 
7      for each vertex  $v \in \text{Adj}[u]$ 
8          do RELAX( $u, v, w$ )
```

Dijkstra's Correctness

Claim

Dijkstra's algorithm terminates with $u.d = \delta(s, u)$ for all $u \in V$.

Proof by Induction

Base case: $|S| = 1$ is trivial.

Inductive hypothesis: Assume true for $|S| = k \geq 1$

- Let v be the next node added to S and let $u-v$ be the chosen edge.
- The shortest $s-u$ path plus (u, v) is an $s-v$ path of length $d(v)$.
- Consider any $s-v$ path P . We'll see that it is no shorter than $d(v)$.
- Let $x-y$ be the first edge in P that leaves S , and let P' be the subpath $s-x$.
- P is already too long as soon as it leaves S :

$$w(P) \geq w(P') + w(x, y) \geq d(x) + w(x, y) \geq d(y) \geq d(v)$$

- The last step of the above is because the loop chose v instead of y as the next vertex to include.

Dijkstra's Algorithm Running Time

Assume we store the $v.d$ values in an array of size V , indexed by v .

- initialization takes $O(V)$ time
- inserting into Q and decreasing the key of an element of Q both take $O(1)$ time (because we stored the keys in an array indexed by v)
- EXTRACT-MIN takes $O(V)$ time
- the **while** loop runs $O(V)$ times, each doing an EXTRACT-MIN, so it's $O(V^2)$ overall
- by aggregate analysis, the **for** loop runs $O(E)$ times, and each RELAX call takes $O(1)$ time
- so the overall running time is $O(V^2 + E) = O(V^2)$

Dijkstra's Running Time (Min-Heap)

If the graph is sparse, we can improve the running time by storing the priority queue in a min-heap

- initialization takes $O(V)$
- the **while** loop runs $O(V)$ times, and **EXTRACT-MIN** takes $O(\lg V)$ time each time it runs giving $O(V \lg V)$ for the loop
- by aggregate analysis, the internal **for** loop runs twice for each edge (once for every entry in an adjacency list), and **RELAX** contains an implicit **DECREASE-KEY**, so the time for the inner loop is $O(E \lg V)$
- so the total running time is $O((V + E) \lg V)$ (or likely $O(E \lg V)$)



Questions
