

EE360C: Algorithms

Dynamic Programming (4/4)

Summer 2019

Department of Electrical and Computer Engineering
University of Texas at Austin

Longest Common Subsequence

Longest Common Subsequence

This is a common problem in biological applications: find the longest common sequence of ACTG in two different pieces of DNA

- useful in finding similar proteins, etc.

Given a sequence $X = \langle x_1, x_2, \dots, x_m \rangle$, a sequence $Z = \langle z_1, z_2, \dots, z_k \rangle$ is a **subsequence** of X if there exists a strictly increasing sequence $\langle i_1, i_2, \dots, i_k \rangle$ of indices of X such that for all $j = 1, 2, \dots, k$, $x_{i_j} = z_j$.

(For example $Z = \langle \text{B}, \text{C}, \text{D}, \text{B} \rangle$ is a subsequence of $X = \langle \text{A}, \text{B}, \text{C}, \text{B}, \text{D}, \text{A}, \text{B} \rangle$.)

Our goal is, given two sequences X and Y , find the longest common subsequence of the two sequences.

Characterizing the Longest Common Subsequence

Theorem: Optimal Substructure of an LCS

Let $X = \langle x_1, x_2, \dots, x_m \rangle$ and $Y = \langle y_1, y_2, \dots, y_n \rangle$ be sequences and let $Z = \langle z_1, z_2, \dots, z_k \rangle$ be any LCS of X and Y .

1. If $x_m = y_n$, then $z_k = x_m = y_n$ and Z_{k-1} is an LCS of X_{m-1} and Y_{n-1} .
2. If $x_m \neq y_n$, then $z_k \neq x_m$ implies that Z is an LCS of X_{m-1} and Y .
3. If $x_m \neq y_n$, then $z_k \neq y_n$ implies that Z is an LCS of X and Y_{n-1} .

Proof (Part I)

If $z_k \neq x_m$, we could append $x_m = y_n$ to Z to get a common subsequence with length $k + 1$, which contradicts the premise. So $z_k = x_m = y_n$. And therefore Z_{k-1} , the prefix of Z with z_k removed is a longest common subsequence of X_{m-1} and Y_{n-1} .

Proof (Parts II and III)

If $z_k \neq x_m$, then Z is a common subsequence of X_{m-1} and Y . If there were a common subsequence W of X_{m-1} and Y longer than k , then W would also be a subsequence of X and Y , which contradicts the premise that Z was an LCS.

A Recursive Solution to LCS

If $x_m = y_n$, there is one subproblem: find the LCS of X_{m-1} and Y_{n-1} . If $x_m \neq y_n$, then there are two subproblems to solve: find the LCS of X_{m-1} and Y and the LCS of X and Y_{n-1} .

$$c[i, j] = \begin{cases} 0 & \text{if } i = 0 \text{ or } j = 0 \\ c[i-1, j-1] + 1 & \text{if } i, j > 0 \text{ and } x_i = y_j \\ \max(c[i, j-1], c[i-1, j]) & \text{if } i, j > 0 \text{ and } x_i \neq y_j \end{cases}$$

Computing the Length of an LCS

We use dynamic programming to compute solutions to the $\Theta(mn)$ subproblems in a bottom-up fashion

- we create a table $c[0 \dots m, 0 \dots n]$ and compute its entries in row major order (filling in the first row, left to right, then the second row, etc.)
- we keep a second table $b[0 \dots m, 0 \dots n]$ whose entries point to the table entry that corresponds to the best subproblem for the problem $b[i, j]$ (to help in reconstructing the optimal solution)

Computing the Length of an LCS (cont.)

LCS-LENGTH(X, Y)

```
1   $m = \text{length}[X]$ 
2   $n = \text{length}[Y]$ 
3  for  $i = 1$  to  $m$ 
4       $c[i, 0] = 0$ 
5  for  $j = 1$  to  $n$ 
6       $c[0, j] = 0$ 
7  for  $i = 1$  to  $m$ 
8      for  $j = 1$  to  $n$ 
9          if  $x_i = y_j$ 
10              $c[i, j] = c[i - 1, j - 1] + 1$ 
11              $b[i, j] = \nwarrow$ 
12         else if  $c[i - 1, j] \geq c[i, j - 1]$ 
13              $c[i, j] = c[i - 1, j]$ 
14              $b[i, j] = \uparrow$ 
15         else  $c[i, j] = c[i, j - 1]$ 
16              $b[i, j] = \leftarrow$ 
17 return  $c$  and  $b$ 
```

An Example

		j	0	1	2	3	4	5	6
		y_j		B	D	C	A	B	A
i	x_i								
0	x_i		0	0	0	0	0	0	0
1	A		0	↑	↑	↑	↖1	←1	↖1
2	B		0	↖1	←1	←1	↑1	↖2	←2
3	C		0	↑1	↑1	↖2	←2	↑2	↑2
4	B		0	↖1	↑1	↑2	↑2	↖3	←3
5	D		0	↑1	↖2	↑2	↑2	↑3	↑3
6	A		0	↑1	↑2	↑2	↖3	↑3	↖4
7	B		0	↖1	↑2	↑2	↑3	↖4	↑4

Shortest Paths


Shortest Paths

Shortest Path Problem

Given a direct graph $G = (V, E)$ with the edge weights c_{vw} (which can include negative edge weights), find the shortest path from node s to node t .

Example

Nodes represent agents in a financial setting and c_{vw} is the cost of a transaction in which we buy from agent v and sell immediately to w



Negative Cost Cycles

Observation

If some path from s to t contains a negative cost cycle, there does not exist a shortest s - t path; otherwise there exists a path that is a simple (i.e. does not repeat nodes) and hence has at most $n - 1$ edges.

Shortest Paths and Dynamic Programming

Definition: $OPT(i, v)$ is the length of the shortest v - t path P using at most i edges.

- Case 1: P uses at most $i - 1$ edges
 - $OPT(i, v) = OPT(i - 1, v)$
- Case 2: P uses exactly i edges
 - if (v, w) is the first edge, the OPT uses (v, w) , and then selects the best w - t path using at most $i - 1$ edges

$$OPT(i, v) = \begin{cases} 0 & \text{if } i = 0 \\ \min(OPT(i - 1, v), \min_{(v, w) \in E} (OPT(i - 1, w) + c_{vw})) & \text{otherwise} \end{cases}$$

Remark

If no negative cycles, then $OPT(n - 1, v)$ = length of shortest v - t path.

Implementation

```
Shortest-Path( $G, t$ ) {  
    foreach node  $v \in V$   
         $M[0, v] \leftarrow \infty$   
     $M[0, t] \leftarrow 0$   
  
    for  $i = 1$  to  $n-1$   
        foreach node  $v \in V$   
             $M[i, v] \leftarrow M[i-1, v]$   
            foreach edge  $(v, w) \in E$   
                 $M[i, v] \leftarrow \min \{ M[i, v], M[i-1, w] + c_{vw} \}$   
}
```

Complexity

$O(mn)$ time, $O(n^2)$ space

Finding shortest path

Maintain a "successor" for each table entry.

Distance Vector Routing

Communication Network

- nodes = routers
- edges = direct communication link
- cost of edge = delay on link (which is naturally non-negative, but we use Bellman-Ford algorithm anyway!)

Dijkstra's Algorithm

Requires global information of network

Bellman-Ford

Uses only local knowledge of neighboring nodes

Synchronization

We don't expect the routers to run in lock-step; Bellman-Ford can tolerate asynchronous routing updates and can be shown to still converge on the correct shortest path.

Distance Vector Protocol

Distance Vector Protocol


- Each router maintains a vector of shortest path lengths to every other node (distances) and the first hop on each path (directions)
- Algorithm: each router performs n separate computations, one for each potential destination node
- “Routing by rumor”

Examples

RIP, Xerox XNS RIP, Novell's IPX RIP, Cisco's IGRP, DEC's DNA Phase IV, AppleTalk's RTMP

Caveat

Edge costs may change during the algorithm; inconsistent router state can lead to the **counting to infinity** problem



Path Vector Protocols

Link State Routing

- Each router stores the entire path (not just the distance and next hop)
- Based on Dijkstra's algorithm
- Avoids the counting to infinity problem
- Requires significantly more storage

Examples

Border Gateway Protocol (BGP), Open Shortest Path First (OSPF)

Questions
